

Answers from your own knowledge — spoken into existence.

The ambient co-pilot that surfaces grounded, cited answers live, mid-conversation.

THE PROBLEM

Customer-facing teams live on calls and drown in documents. Sellers spend only ~25% of their time selling — ~64% goes to hunting for information. When a question lands mid-call, the answer is in a doc somewhere, but finding it takes 30 seconds the conversation doesn't have. Today's live-assist tools guess from a general model with no access to your real knowledge: confident, often wrong.

THE SOLUTION

Relay listens to a live conversation and, the instant a topic comes up, retrieves the answer from your own documents — sub-10 ms via Moss — and renders a grounded, **cited** card on screen in under half a second. Not a chatbot. Grounded retrieval, in real time.

WHY NOW

Voice models got cheap and fast — retrieval became the bottleneck. Sub-10 ms retrieval just made real-time grounding possible. The viral live-assist tools proved the demand but have no company-knowledge grounding. That gap is our entire product.

PRODUCT — ONE ENGINE, THREE AGENTS

Relay Live · Co-Pilot

Surfaces context on live calls. Never interrupts.

Relay Desk · Support

Resolves with the right doc + customer history.

Relay Intake · Lead Gen

Qualifies and routes inbound leads in conversation.

<10ms

knowledge
retrieval (Moss)

<500ms

spoken question to
card

100%

answers cited to
a source

MARKET & MOAT

Customer-facing teams at 10–50 person companies — underserved between DIY tools and six-figure enterprise platforms. Real-time retrieval at conversation latency is hard, and citation-grounding builds trust generic assistants can't. Every call makes the knowledge base and memory better — a compounding data advantage.

THE ASK

Feedback, design partners, and a conversation about replacing the scramble for information at scale. Waitlist live within 24h — first design partners from the teams who feel the pain most.